



Development of a Sky Condition Classifier Using All-Sky Images

Mingyu Jeon, Rizchel Masong, Yi-Chieh Chang



**meme explainer: the sky when there is a cool astronomical event.*

A3N 2025 (18-22 Aug 2025)
KIAS, Seoul, South Korea.

Motivation

DESI

Higashi-Hiroshima Observatory



Photo: <https://dive-hiroshima.com/en/explore/1685/>




Photo: <https://www.enkey.it/en/2019/11/20/desi-the-super-telescope-that-will-help-to-draw-a-new-map-of-the-sky/>

Motivation

Our data

Higashi-Hiroshima Observatory



Photo: <https://dive-hiroshima.com/en/explore/1685/>

DESI



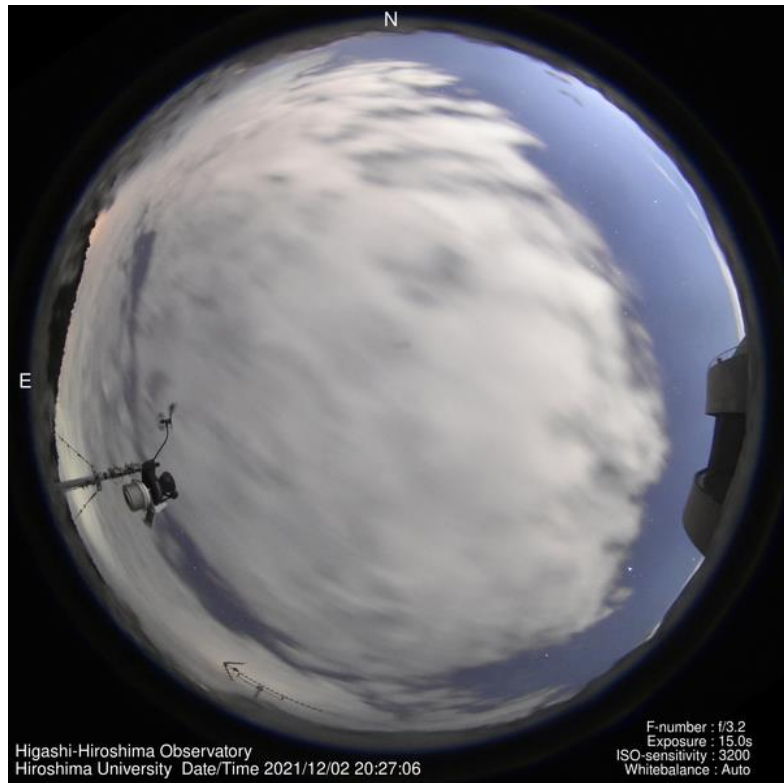




Photo: <https://www.enkey.it/en/2019/11/20/desi-the-super-telescope-that-will-help-to-draw-a-new-map-of-the-sky/>


Motivation



Motivation



Motivation



Environment Monitor at Higashi-Hiroshima Observatory



Current Time (This PC)
2025/08/21 (Thu) 12:27:28

Local Sidereal Time
-

METEOROGRAM						
	Time	Temperature (deg's C)	Humidity (%)	Wind Dir.	Wind Sp. (m/s)	Rain (mm/5min)
Outside	12:10	+28.1	77	SSE	2.2	0.0
Dome Inside	12:05	+22.5	47	-	-	-

RAINDROP SENSOR (at 12:31:04)				TELESCOPE Temp (at 12:24:00)		
North		South		TopRingS	CenSectS	PillMidE
Level 0	DRY	Level 0	DRY	+25.2	+22.6	+23.6

TELESCOPE (at 12:16:02)			
Control PC	Position	GRB mode	GRB obs status
OFF	Azimuth - Altitude -	OFF	-


DOME		
DomeSlit	MirrorCover	Azimuth
CLOSED	CLOSED	+0.0 (-)

Link

- [Meteogram of HHO](#)
- [Hiroshima Astrophysical Science Center](#)
- Observatory monitoring camera
 - [\[Rikubetsu\(Ginganomori\)\]](#) [\[Nayoro\(Hokkaido Univ.\)\]](#) [\[Gunma\(GAO\)\]](#) [\[Saitama Univ.\]](#)
 - [\[Akeno\(Tokyo-kogyo Univ.\)\]](#) [\[Kyoto\(KAO\)\]](#) [\[Osaka\(Osaka Kyoiku Univ.\)\]](#)
 - [\[Sayo\(NHAO\)\]](#) [\[Okayama\(OAO\)\]](#) [\[Bisei\(JSGA\)\]](#) [\[Iriki\(Kagoshima Univ.\)\]](#)
 - [\[Sutherland\(SuperWASP.group\)\]](#)


Method

- Data : 4,984 (Cloudy) + 3,822 (Fine) = 8,806 Images
- Binary Classification
- CNN



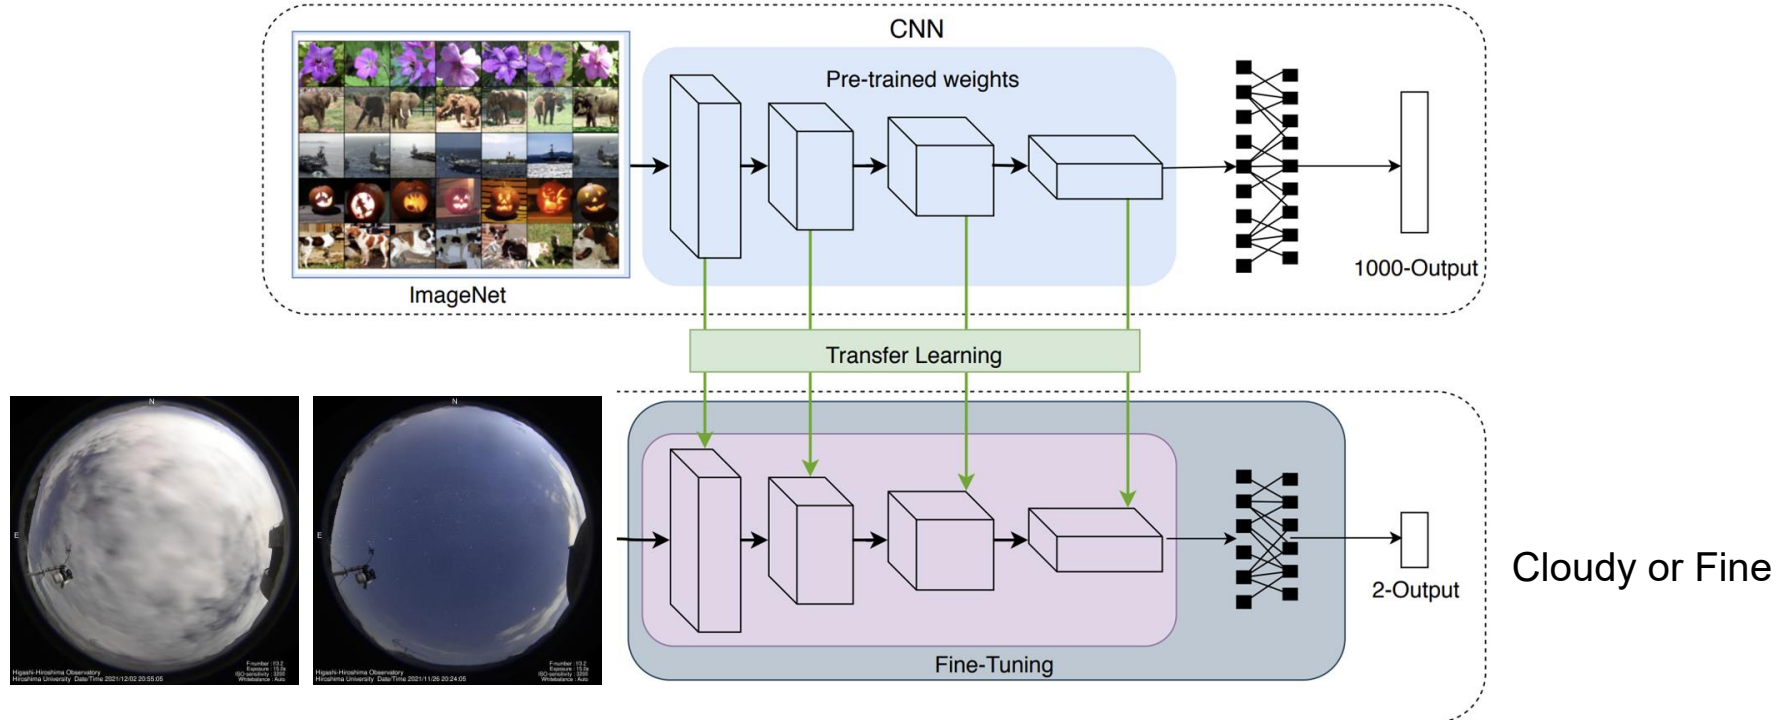
Method

- Data : 4,984 (Cloudy) + 3,822 (Fine) = 8,806 Images
- Binary Classification
- CNN = Feature Extractor + Classifier



Transfer Learning

Idea: Pre-trained CNN layers can extract useful features from (almost) any image



A PyTorch library for pre-trained vision models




Hugging Face

- Aggregating Nested Transformers - <https://arxiv.org/abs/2105.12723>
- BEiT - <https://arxiv.org/abs/2106.08254>
- BEiT-V2 - <https://arxiv.org/abs/2208.06366>
- BEiT3 - <https://arxiv.org/abs/2208.10442>
- Big Transfer ResNetV2 (BiT) - <https://arxiv.org/abs/1912.11370>
- Bottleneck Transformers - <https://arxiv.org/abs/2101.11605>
- CaiT (Class-Attention in Image Transformers) - <https://arxiv.org/abs/2103.17239>
- CoaT (Co-Scale Conv-Attentional Image Transformers) - <https://arxiv.org/abs/2106.04803>
- CoAtNet (Convolution and Attention) - <https://arxiv.org/abs/2106.04803>
- ConvNeXt - <https://arxiv.org/abs/2201.03545>
- ConvNeXt-V2 - <https://arxiv.org/abs/2301.00808>
- ConViT (Soft Convolutional Inductive Biases Vision Transformers) - <https://arxiv.org/abs/2106.04803>
- CspNet (Cross-Stage Partial Networks) - <https://arxiv.org/abs/1911.11929>
- DeiT - <https://arxiv.org/abs/2012.12877>
- DeiT-III - <https://arxiv.org/pdf/2204.07118.pdf>
- DenseNet - <https://arxiv.org/abs/1608.06993>
- DLA - <https://arxiv.org/abs/1707.06484>
- DPN (Dual-Path Network) - <https://arxiv.org/abs/1707.01629>
- EdgeNeXt - <https://arxiv.org/abs/2206.10589>
- EfficientFormer - <https://arxiv.org/abs/2206.01191>
- EfficientFormer-V2 - <https://arxiv.org/abs/2212.08059>
- EfficientNet (MBConvNet Family)
 - EfficientNet NoisyStudent (B0-B7, L2) - <https://arxiv.org/abs/1911.04252>
 - EfficientNet AdvProp (B0-B8) - <https://arxiv.org/abs/1911.09665>
 - EfficientNet (B0-B7) - <https://arxiv.org/abs/1905.11946>
 - EfficientNet-EdgeTPU (S, M, L) - <https://ai.googleblog.com/2019/08/efficientnet-edge-tpu.html>

Model

Credit: EfficientNet paper (Tan & Le, 2019 arXiv)





We use 3 models with ~20M params

EfficientNet-B4	17.6M
DenseNet-201	18.1M
ResNet-34	21.3M

✂ params of only CNN layers


Similar #parameters !

Model Explanation: ResNet & DenseNet



\oplus : Element-wise addition

ResNet-34
= ResNet with 34 layers


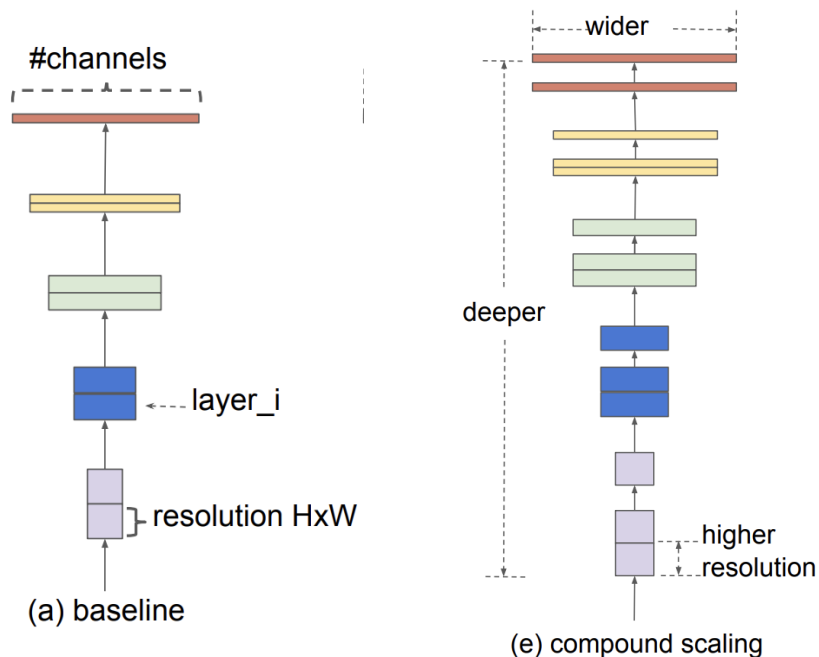


\odot : Channel-wise concatenation

DenseNet-201
= DenseNet with 201 layers

Model Explanation: EfficientNet

- (1) Find a baseline architecture using Neural Architecture Search
- (2) Scale up the baseline architecture using Compound Scaling



EfficientNet-B4

= EfficientNet with compound scaling stage 4

$$\text{depth: } d = \alpha^\phi$$

$$\text{width: } w = \beta^\phi$$


$$\text{resolution: } r = \gamma^\phi$$

$$\text{s.t. } \alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$$

$$\alpha \geq 1, \beta \geq 1, \gamma \geq 1$$

Results (1): Finetune > Scratch > Frozen


Scratch	Train the model from random initialization
Frozen	Freeze feature extractor layers of a pre-trained model and only train the final classifier
Finetune	Start with a pre-trained model and update all layers on the new dataset



Results (1): Finetune > Scratch > Frozen


Why 'Frozen' is worst?

- Train only a linear classifier (output = $W * \text{feature vector} + b$)
- Test if features extracted from pre-trained CNN are linearly separable for our data
- Perform worst since features extracted from pre-trained CNN are not necessarily linearly separable for our data




Results (2): EfficientNet > DenseNet > ResNet

Model (# params)	Test Accuracy (best epoch) best epoch = best validation accuracy		
	scratch	finetune	frozen
EfficientNet-B4 (17.6M)	1.0000	0.9989	0.9524
DenseNet-201 (18.1M)	0.9966	0.9977	0.9705
ResNet-34 (21.3M)	0.9966	0.9977	0.9388



Results (2): EfficientNet > DenseNet > ResNet



EfficientNet-B4	17.6M
DenseNet-201	18.1M
ResNet-34	21.3M

✱ params of only CNN layers

Explainable AI: CAM-based methods

Advanced AI explainability for PyTorch

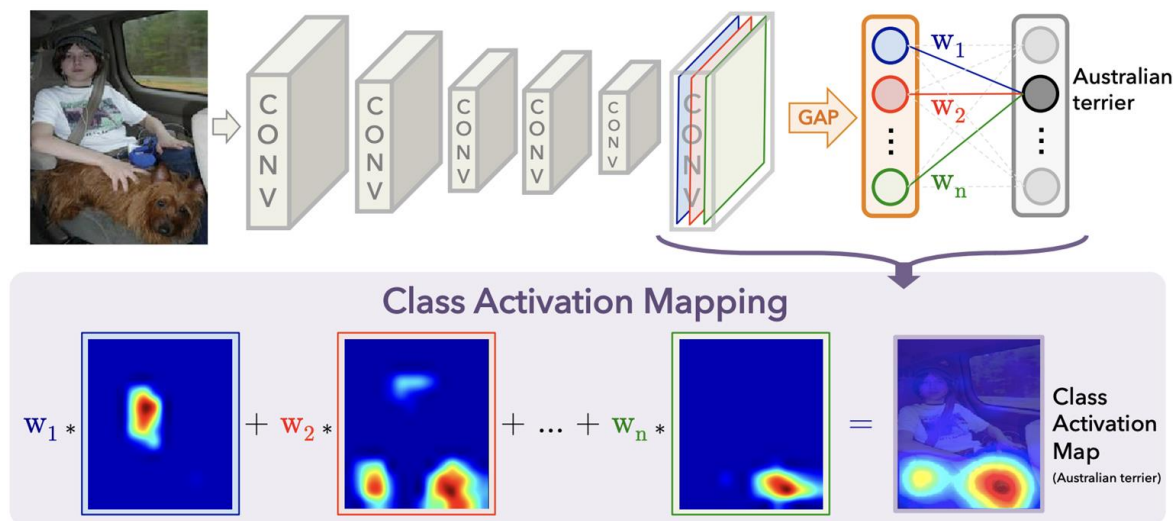
`pip install grad-cam`

Documentation with advanced tutorials: <https://jacobgil.github.io/pytorch-gradcam-book>

Class Activation Map (CAM)

Idea: If a CNN classifier works well, the last convolutional feature maps usually contain class-specific spatial information.

→ CAM-based methods (e.g., Grad-CAM, Eigen-CAM, etc.) visualize this information in various ways.



Method
GradCAM
HiResCAM
GradCAMElementWise
GradCAM++
XGradCAM
AblationCAM
ScoreCAM
EigenCAM
EigenGradCAM
LayerCAM
FullGrad
Deep Feature Factorizations
KPCA-CAM

Grad-CAM (Gradient-weighted Class Activation Map)

Grad-CAM for a class c

ReLU \rightarrow only positive influence... pixels whose intensity is increased to increase score y^c


$$L_{\text{Grad-CAM}}^c = \text{ReLU} \left(\underbrace{\sum_k \alpha_k^c A^k}_{\text{linear combination}} \right)$$

α_k^c is the “importance” of feature map k , A^k , for the class c


$$\alpha_k^c = \overbrace{\frac{1}{Z} \sum_i \sum_j}^{\text{global average pooling}} \underbrace{\frac{\partial y^c}{\partial A_{ij}^k}}_{\text{gradients via backprop}}$$

Red pixels = higher contribution to the model's prediction
Blue pixels = lower contribution to the model's prediction

Actual / Prediction / Probability
EfficientNet-B4 | cloudy / cloudy / 100%



cloudy
Grad-CAM



Weakly-supervised semantic segmentation

Can an image classifier trained only with image-level labels produce pixel-level labels?



DenseNet-201

(finetune, last epoch)

Grad-CAM > threshold



×


Circle mask



Grad-CAM

(Grad-CAM > thrs)*circle

Actual / Prediction / Probability
DenseNet-201 | cloudy / cloudy / 100%



EfficientNet-B4


(finetune, last epoch)

For this image, EfficientNet-B4 detects more cloud regions (via Grad-CAM) than DenseNet-201.


Actual / Prediction / Probability
EfficientNet-B4 | cloudy / cloudy / 100%



Grad-CAM



(Grad-CAM > thrs)*circle



ResNet-34


(finetune, last epoch)

ResNet-34 uses different regions than DenseNet-201 and EfficientNet-B4 to classify the sky as cloudy

Actual / Prediction / Probability
ResNet-34 | cloudy / cloudy / 100%



Grad-CAM



(Grad-CAM > thr)*circle



Summary and Future works

- Summary
 - CNN Architecture: EfficientNet > DenseNet > ResNet
 - Training Strategy: Fine-Tuning > From-Scratch > Frozen
 - Grad-CAM provides interpretability for CNN predictions
 - Grad-CAM results can be used for weakly-supervised semantic segmentation
- Future Works
 - Image augmentation (rotation, brightness, contrast, etc.)
 - Dataset expansion (different seasons, times of day)
 - Multi-class classification (e.g., partly cloudy, rainy)
 - Real-time processing (end-to-end pipeline)
 - Robustness testing (e.g., moonlight conditions)
 - Time-series prediction

Appendix

Binary Classification

Task: Data $x \rightarrow$ Two Discrete Labels $y \in \{0, 1\}$

1: cloudy

0: fine


Assume the labels follow a Bernoulli distribution

$$p(y|\lambda) = (1 - \lambda)^{1-y} \lambda^y \quad ; \quad \lambda = (\text{probability that } y=1)$$

Train a model $f_\theta(x)$ to predict the distribution parameter λ

$$\lambda = \sigma(f_\theta(x)) \quad ; \quad \sigma = \text{sigmoid function from } \mathbb{R} \text{ to } [0, 1]$$

sigmoid function



The loss function (**binary cross entropy**) is the negative log-likelihood of the training set

$$\begin{aligned} L(\theta) &= - \sum_{i=1}^N \log p(y_i | \sigma(f_\theta(x_i))) \\ &= - \sum_{i=1}^N (1 - y_i) \log[(1 - \sigma(f_\theta(x_i)))] + y_i \log[\sigma(f_\theta(x_i))] \end{aligned}$$

Detail

- Dataset
 - Total: 8,806 images (4,984 Cloudy + 3,822 Fine) 600x600 size
 - Split (8:1:1):
 - Cloudy: 3,987 / 498 / 499
 - Fine: 3,057 / 382 / 383
 - Total: 7,044 / 880 / 882 (Train / Val / Test)
 - Preprocessing: Resized and normalized using the pre-trained model's pipeline
- Training Setup
 - Epochs: 20
 - Batch size: 64
 - Learning rate: $2e-4$

Pre-trained Models: ResNet-34

```
Compose(  
    Resize(size=235, interpolation=bicubic, max_size=None, antialias=True)  
    CenterCrop(size=(224, 224))  
    MaybeToTensor()  
    Normalize(mean=tensor([0.4850, 0.4560, 0.4060]), std=tensor([0.2290, 0.2240, 0.2250]))  
)
```

timm / **resnet34.a1_in1k**   like 0 Follow — PyTorch Image Models 537

 Image Classification  timm  PyTorch  Safetensors  Transformers  arxiv:21

 **Model card**  Files and versions  Community



Model card for resnet34.a1_in1k

A ResNet-B image classification model.


This model features:






- ReLU activations
- single layer 7x7 convolution with pooling
- 1x1 convolution shortcut downsample





Trained on ImageNet-1k in `timm` using recipe template described below.

Pre-trained Models: DenseNet-201

```
Compose(  
  Resize(size=256, interpolation=bicubic, max_size=None, antialias=True)  
  CenterCrop(size=(224, 224))  
  MaybeToTensor()  
  Normalize(mean=tensor([0.4850, 0.4560, 0.4060]), std=tensor([0.2290, 0.2240, 0.2250]))  
)
```

timm **timm/densenet201.tv_in1k**  ♡ like 0 Follow timm PyTorch Image Models 537

 Image Classification timm  PyTorch  Safetensors  Transformers  imagenet-1k



 **Model card**  Files and versions  Community  Edit


Model card for densenet201.tv_in1k







A DenseNet image classification model. Trained on ImageNet-1k (original torchvision weights).




Pre-trained Models: EfficientNet-B4


```
Compose(  
    Resize(size=365, interpolation=bicubic, max_size=None, antialias=True)  
    CenterCrop(size=(320, 320))  
    MaybeToTensor()  
    Normalize(mean=tensor([0.4850, 0.4560, 0.4060]), std=tensor([0.2290, 0.2240, 0.2250]))  
)
```

 **timm/efficientnet_b4.ra2_in1k** 

 like 0 Follow PyTorch Image Models 5

 Image Classification timm  PyTorch  Safetensors  Transformers  imagenet-1k 

 **Model card**  Files and versions  Community

 Edit mode

Model card for efficientnet_b4.ra2_in1k


A EfficientNet image classification model. Trained on ImageNet-1k in `timm` using recipe template described below.

Red pixels = higher contribution to the model's prediction
Blue pixels = lower contribution to the model's prediction

Actual / Prediction / Probability
EfficientNet-B4 | fine / fine / 100%



fine
Grad-CAM



Wrong prediction (TODO)